

B. K. Birla College (Autonomous), Kalyan (Department of Information Technology)



This is to certify that Mrs. **Ruchita S. Kamble**, Roll No. **15** ,
Exam Seat No. **23258715** has satisfactorily completed the Practical in
BIG DATA ANALYTICS of Semester **II** for the fulfilment of M.SC.- IT(Cloud
Computing) for the year **2024-2025**.

Signature of Examiners

Date:- 03 June 2025

INDEX

No.	Practical Name	Date	Page No.
1	Exploring Big Data Characteristics with Real Datasets	20-May-2025	2
2	Comparative Analysis of Traditional BI vs Big Data Tools	20-May-2025	6
3	Implement K-Means and DBSCAN Clustering	20-May-2025	8
4	Association Rules Mining using Apriori Algorithm	21-May-2025	12
5	Regression Analysis and Evaluation	21-May-2025	14
6	Building and Evaluation of Classification Model	21-May-2025	17
7	Time Series Forecasting using ARIMA	21-May-2025	20
8	Text Analysis and Sentiment Detection	22-May-2025	23
9	Hadoop HDFS and MapReduce Word Count	22-May-2025	27
10	Building a Spark Application with PySpark	22-May-2025	29

Practical 1: Exploring Big data Characteristics with Real Datasets

#load the dataset

```
[1] import pandas as pd
df=pd.read_csv("OnlineRetail.csv", encoding='ISO-8859-1') #encoding- to handle special characters in the dataset.
df.head()
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	12/1/2010 8:26	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	12/1/2010 8:26	2.75	17850.0	United Kingdom
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	12/1/2010 8:26	3.39	17850.0	United Kingdom

```
[2 ] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   InvoiceNo    541909 non-null object
1   StockCode   541909 non-null object
2   Description  540455 non-null object
3   Quantity    541909 non-null int64
4   InvoiceDate  541909 non-null object
5   UnitPrice   541909 non-null float64
6   CustomerID  406829 non-null float64
7   Country     541909 non-null object
dtypes: float64(2), int64(1), object(5)
memory usage: 33.1+ MB
```

Volume -- Size of Data

```
[3] print("Shape:", df.shape)
#shows the size: (rows, columns) → indicates how much data you have.
```

Shape: (541909, 8)

```
[4 ] print("Memoery Usage (MB):",df.memory_usage(deep=True).sum()/1024**2)
#calculates how much memory the DataFrame is consuming, in megabytes (MB)
```

Memoery Usage (MB): 173.12356662750244

#Velocity --Speed of Data Generation & Processing

```
[ ] df['InvoiceDate']=pd.to_datetime(df['InvoiceDate'],dayfirst=True, errors='coerce')
#Converts the InvoiceDate column from string to datetime.
#dayfirst=True handles dates in DD/MM/YYYY format.
#errors='coerce' converts invalid dates to NaT (missing date).
```

```
[5 ] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   InvoiceNo    541909 non-null object
1   StockCode   541909 non-null object
2   Description  540455 non-null object
3   Quantity    541909 non-null int64
4   InvoiceDate  232959 non-null datetime64[ns]
5   UnitPrice   541909 non-null float64
6   CustomerID  406829 non-null float64
7   Country     541909 non-null object
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB
```

#Variety --Different Data Types and Sources

```
[6 ] print("Data Type:\n", df.dtypes)
#df.dtypes shows the data type of each column (e.g., object, float64, datetime64)
```

```
Data Type:
InvoiceNo      object
StockCode      object
Description     object
Quantity       int64
InvoiceDate    datetime64[ns]
UnitPrice      float64
CustomerID     float64
Country        object
dtype: object
```

```
[7 ] print("Variety Count:\n", df.dtypes.value_counts())
#.value_counts() on dtypes tells how many columns belong to each type.
```

```
Variety Count:
object      4
float64     2
int64       1
datetime64[ns]  1
Name: count, dtype: int64
```

#Veracity-- Data Quality and Trustworthiness

```
[8 ] print("Missing Values:\n", df.isnull().sum())
#isnull().sum() counts how many missing values are in each column.
```

Missing Values:

```
InvoiceNo      0
StockCode      0
Description    1454
Quantity       0
InvoiceDate    308950
UnitPrice      0
CustomerID     135080
Country        0
dtype: int64
```

```
[9 ] print("Duplicate Records:\n", df.duplicated().sum())
#duplicated().sum() counts how many duplicate rows exist.
```

Duplicate Records:
5269

```
# Cleaning Steps
#Removes rows where CustomerID or InvoiceDate is missing.
#This improves data quality (veracity).
[10 ] df_clean= df.dropna(subset=['CustomerID'])
```

```
[11] print("After cleaning:", df_clean.shape)
```

After cleaning: (406829, 8)

```
[12 ] print("Missing Values:\n", df_clean.isnull().sum())
```

Missing Values:

```
InvoiceNo      0
StockCode      0
Description     0
Quantity       0
InvoiceDate    234047
UnitPrice      0
CustomerID      0
Country        0
dtype: int64
```

```
[13] df_clean1= df.dropna(subset=['InvoiceDate'])
```

```
[14 ] print("Missing Values:\n", df_clean1.isnull().sum())
```

Missing Values:

```
InvoiceNo      0
StockCode      0
Description    658
Quantity       0
InvoiceDate     0
UnitPrice      0
CustomerID     60177
Country        0
```

#Calculate Revenue per Transaction

```
[15 ] df['TotalPrice']=df['Quantity'] * df['UnitPrice']
```

#Creates a new column TotalPrice, representing revenue per line item.

```
[16 ] print("Top Products:\n", df['Description'].value_counts().head())
```

#Finds the most frequently sold products by counting values in the Description column.

Top Products:

```
Description
WHITE HANGING HEART T-LIGHT HOLDER  2369
REGENCY CAKESTAND 3 TIER            2200
JUMBO BAG RED RETROSPOT              2159
PARTY BUNTING                       1727
LUNCH BAG RED RETROSPOT              1638
Name: count, dtype: int64
```

```
[ 17 ] print("Top Countries by Revenue:\n",
df.groupby('Country')['TotalPrice'].sum().sort_values(ascending=False).head())
```

Top Countries by Revenue:

```
Country
United Kingdom  8187806.364
Netherlands    284661.540
EIRE            263276.820
Germany        221698.210
France         197403.900
Name: TotalPrice, dtype: float64
```

#Groups the data by Country.
#Sums the TotalPrice for each country.
#Sorts countries by revenue in descending order.
#Shows the top 5 countries with the highest sales revenue.

```
[ 18 ] df['TotalPrice'] = df['Quantity'] * df['UnitPrice']
```

```
[19 ] print("Top Products:\n", df['Description'].value_counts().head())
```

Top Products:

```
Description
WHITE HANGING HEART T-LIGHT HOLDER  2369
REGENCY CAKESTAND 3 TIER            2200
JUMBO BAG RED RETROSPOT              2159
PARTY BUNTING                       1727
LUNCH BAG RED RETROSPOT              1638
Name: count, dtype: int64
```

Practical 2: Comparative Analysis of Traditional BI vs Big Data Tools

```
[1] #comapre analysis and data visulization
import pandas as pd
df = pd.read_csv("OnlineRetail.csv", encoding='ISO-8859-1')
```

```
[2] df.head()
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	12/1/2010 8:26	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	12/1/2010 8:26	2.75	17850.0	United Kingdom
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	12/1/2010 8:26	3.39	17850.0	United Kingdom

```
[3]df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'],dayfirst=True,errors='coerce')
```

```
[4] # compute total price
df['TotalPrice'] = df['Quantity'] * df['UnitPrice']
```

```
[5]# group by the month of the InvoiceDate
monthly_sales = df.groupby(df['InvoiceDate'].dt.to_period('M')).agg( {
    'Quantity': 'sum',
    'TotalPrice': 'sum'
}).reset_index()
```

```
print(monthly_sales.head())
```


```
InvoiceDate  Quantity  TotalPrice
0  2010-01      26814    58635.56
1  2010-02      21023    46207.28
2  2010-03      14830    45620.46
3  2010-05      16395    31383.95
4  2010-06      21419    53860.18
```

```
[6] # 1. Reset index so 'InvoiceDate' is a column
monthly_sales = monthly_sales.reset_index()
```

```
# 2. Convert the PeriodIndex column to Timestamps
monthly_sales['InvoiceDate'] = monthly_sales['InvoiceDate'].dt.to_timestamp()
```

```
# 3. Plot
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(14,6))
plt.plot(
    monthly_sales['InvoiceDate'],
    monthly_sales['TotalPrice'],
    marker='o',
    label='Total Sales'
)
plt.plot(
    monthly_sales['InvoiceDate'],
    monthly_sales['Quantity'],
    marker='o',
    label='Total Sales'
)
plt.title('Monthly Total Sales')
plt.xlabel('Month')
plt.ylabel('TotalPrice')
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.show()
```



Practical 3: Implement K-Means and DBSCAN Clustering

```
[1] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans,DBSCAN
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import silhouette_score
```

```
[2] df = pd.read_csv('Mall_Customers.csv')
print(df.head())
```

CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1 Male	19	15	39
1	2 Male	21	15	81
2	3 Female	20	16	6
3	4 Female	23	16	77
4	5 Female	31	17	40

```
[3] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                200 non-null   object
2   Age                  200 non-null   int64
3   Annual Income (k$)    200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
[4] df.describe()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

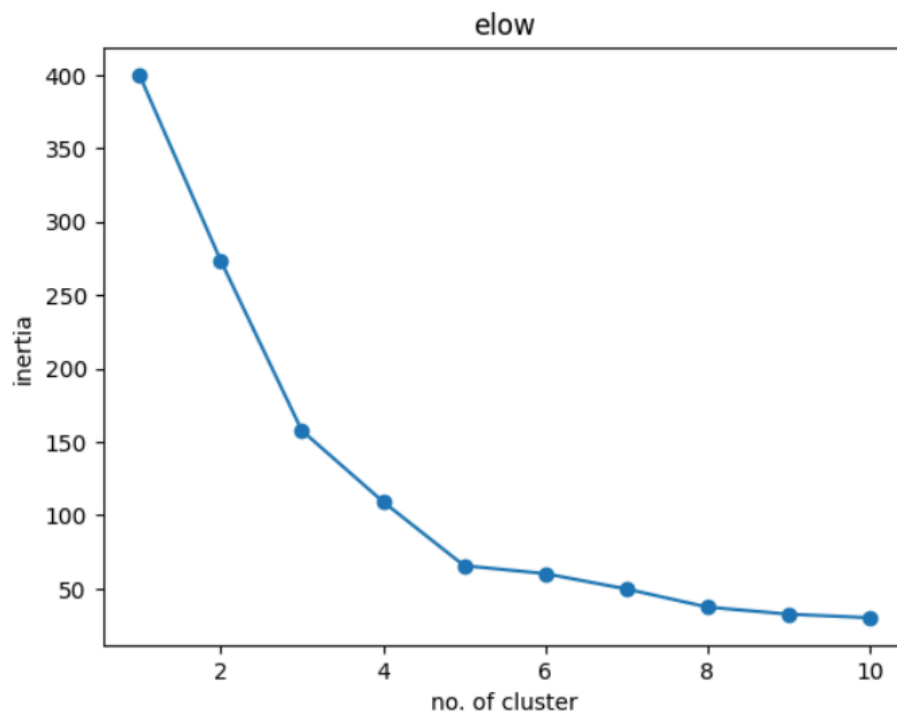
```
[5] df.isnull().sum()
```

```
CustomerID      0  
Gender           0  
Age             0  
Annual Income (k$)  0  
Spending Score (1-100)  0  
dtype: int64
```

```
[6] X = df[['Annual Income (k$)', 'Spending Score (1-100)']]  
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```

```
[7] inertia = []  
for k in range(1,11):  
    kmeans = KMeans(n_clusters=k, random_state=42)  
    kmeans.fit(X_scaled)  
    inertia.append(kmeans.inertia_)
```

```
plt.plot(range(1,11), inertia, marker='o')  
plt.xlabel('no. of cluster')  
plt.ylabel('inertia')  
plt.title('elbow')  
plt.show()
```




```
[8] from sklearn.cluster import KMeans  
import matplotlib.pyplot as plt  
import seaborn as sns
```

import pandas as pd

[9] # Assuming X_scaled is a scaled version of your input data (e.g., with StandardScaler or MinMaxScaler)
and k is defined

```
kmeans = KMeans(n_clusters=k, random_state=42)
df['KMeans_cluster'] = kmeans.fit_predict(X_scaled)
plt.figure(figsize=(8,6))
# Correct indexing for x and y: use column indices or names
sns.scatterplot(
    x=X_scaled[:, 0], # 1st feature (e.g., Annual Income)
    y=X_scaled[:, 1], # 2nd feature (e.g., Spending Score)
    hue=df['KMeans_cluster'],
    palette='viridis'
)
```

```
plt.xlabel('Annual Income (scaled)')
plt.ylabel('Spending Score (scaled)')
plt.title('K-Means Clustering')
plt.show()
```



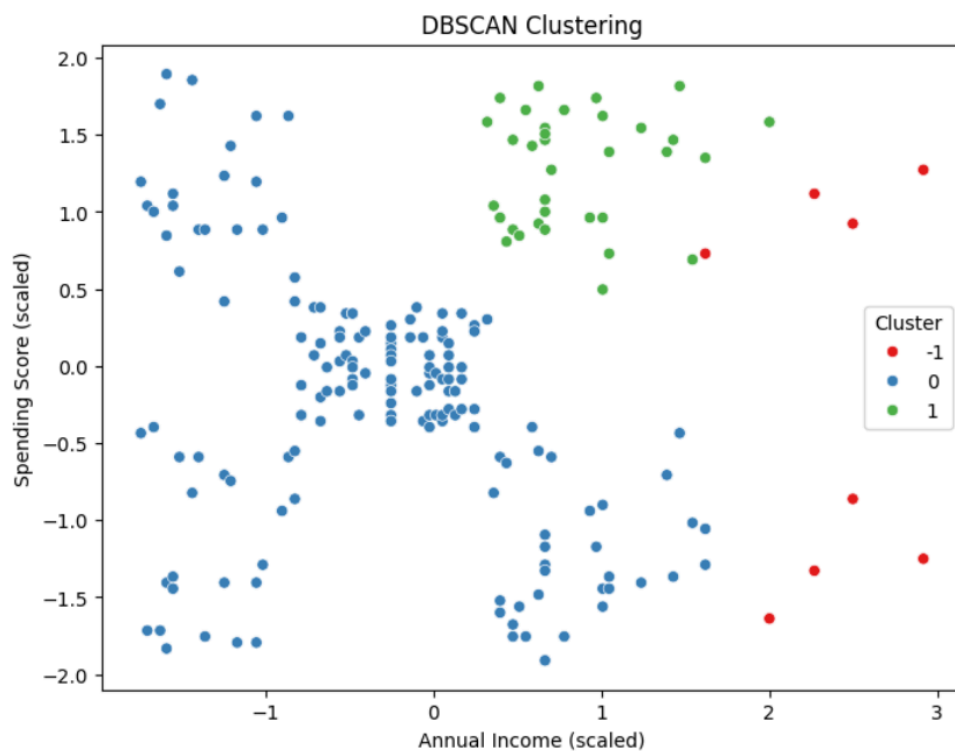
[10] from sklearn.cluster import DBSCAN

```
# Create and fit the DBSCAN model
dbscan = DBSCAN(eps=0.5, min_samples=5) # Corrected: min_samples, not min_sample
df['DBSCAN_Cluster'] = dbscan.fit_predict(X_scaled)
```

```
# In[27]:
plt.figure(figsize=(8,6))
```

```
sns.scatterplot(  
    x=X_scaled[:, 0],  
    y=X_scaled[:, 1],  
    hue=df['DBSCAN_Cluster'],  
    palette='Set1'  
)
```

```
plt.xlabel('Annual Income (scaled)')  
plt.ylabel('Spending Score (scaled)')  
plt.title('DBSCAN Clustering')  
plt.legend(title='Cluster')  
plt.show()
```



Practical 4: Association Rules Mining using Apriori Algorithm

```
[1] get_ipython().system('pip install mlxtend')

[2] import pandas as pd
from mlxtend.preprocessing import TransactionEncoder
from mlxtend.frequent_patterns import apriori, association_rules

[3] transactions=[
    ['Milk','Bread','Butter'],
    ['Bread','Butter'],
    ['Milk','Diaper','Beer','Egg'],
    ['Bread','Butter','Diaper'],
    ['Milk','Diaper','Bread','Butter'],
    ['Beer','Diaper'],
    ['Milk','Butter','Diaper','Egg'],
]

[4] te=TransactionEncoder()
te_ary =te.fit(transactions).transform(transactions)
df_onehot =pd.DataFrame(te_ary, columns=te.columns_)
print("One-hot encoded dataset:")
display(df_onehot)
```

One-hot encoded dataset:

	Beer	Bread	Butter	Diaper	Egg	Milk
0	False	True	True	False	False	True
1	False	True	True	False	False	False
2	True	False	False	True	True	True
3	False	True	True	True	False	False
4	False	True	True	True	False	True
5	True	False	False	True	False	False
6	False	False	True	True	True	True

```
[5] frequent_itemset =apriori(df_onehot, min_support=0.4, use_colnames=True)
print("frequent_itemset")
display(frequent_itemset)
```

frequent_itemset		
	support	itemsets
0	0.571429	(Bread)
1	0.714286	(Butter)
2	0.714286	(Diaper)
3	0.571429	(Milk)
4	0.571429	(Butter, Bread)
5	0.428571	(Butter, Diaper)
6	0.428571	(Butter, Milk)
7	0.428571	(Milk, Diaper)

```
[6] rules = association_rules(frequent_itemset, metric="confidence", min_threshold=0.7)
print(association_rules)
display(rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']])
```

<function association_rules at 0x0000017170A9EC00>

	antecedents	consequents	support	confidence	lift
0	(Butter)	(Bread)	0.571429	0.80	1.40
1	(Bread)	(Butter)	0.571429	1.00	1.40
2	(Milk)	(Butter)	0.428571	0.75	1.05
3	(Milk)	(Diaper)	0.428571	0.75	1.05

Practical 5: Regression Analysis And Evaluation

```
[1] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split

[2] np.random.seed(42)
data= pd.DataFrame({
    'Advertising_spend': np.random.uniform(1000, 5000, 100),
    'Store_size': np.random.uniform(500, 1500, 100),
    'Sales': np.random.uniform(20000, 100000, 100),
})
x=data[['Advertising_spend','Store_size']]
y=data['Sales']

x_train, x_test, y_train, y_test =train_test_split(x, y, test_size=0.2, random_state=42)

model = LinearRegression()
model.fit(x_train, y_train)

y_pred = model.predict(x_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))


r2= r2_score(y_test, y_pred)

print(f'Linear Regression -RMSE: {rmse:.2f}')
print(f'Linear Regression -R^2 score: {r2:.2f}')

plt.scatter(y_test, y_pred)

plt.xlabel("Actual sales")
plt.ylabel("Predicted sales")
plt.title("Actual vs predicted sales")
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--')
plt.show()
```

Linear Regression -RMSE: 26922.70
Linear Regression -R² score: -0.27



```
[3] import pandas as pd
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

# Generating synthetic churn data
churn_data = pd.DataFrame({
    'Monthly_charges': np.random.uniform(20, 120, 100),
    'Tenure': np.random.randint(1, 72, 100),
    'Churn': np.random.choice([0, 1], 100) # 0 = no churn, 1 = churn
})

# Feature matrix and target vector
x = churn_data[['Monthly_charges', 'Tenure']]
y = churn_data['Churn']

# Splitting data into train and test sets
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

# Logistic Regression model
log_model = LogisticRegression()
log_model.fit(x_train, y_train)

# Making predictions
```



```
y_pred = log_model.predict(x_test)
```

```
# Evaluating the model
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
print(f'Logistic Regression - Accuracy: {accuracy:.2f}')
```

```
print("Confusion Matrix:")
```

```
print(conf_matrix)
```

```
print("\nClassification Report:")
```

```
print(classification_report(y_test, y_pred))
```

```
# Visualization (optional but illustrative)
```

```
plt.figure(figsize=(6, 4))
```

```
plt.scatter(y_test, y_pred, alpha=0.6)
```

```
plt.xlabel("Actual Churn")
```

```
plt.ylabel("Predicted Churn")
```

```
plt.title("Actual vs Predicted Churn")
```

```
plt.plot([0, 1], [0, 1], 'r--') # Diagonal reference line
```

```
plt.grid(True)
```

```
plt.show()
```

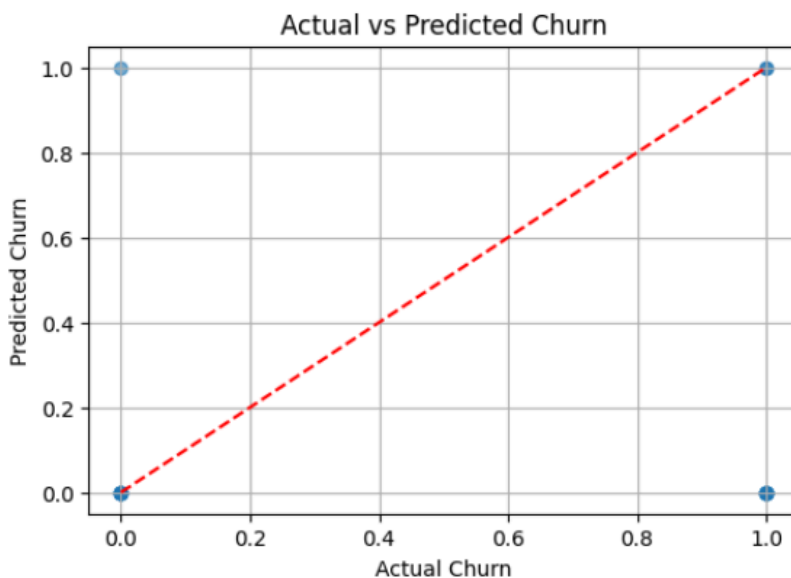
```
Logistic Regression - Accuracy: 0.50
```

```
Confusion Matrix:
```

```
[[ 8  1]
 [ 9  2]]
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.47	0.89	0.62	9
1	0.67	0.18	0.29	11
accuracy			0.50	20
macro avg	0.57	0.54	0.45	20
weighted avg	0.58	0.50	0.43	20



Practical 6: Building and evaluation Classification model

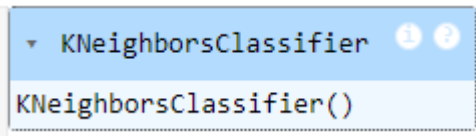
```
[1] import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix

[2] iris=load_iris()##load data
X= iris.data
y=iris.target

[3] X_train, X_test, y_train, y_test =train_test_split(X, y, test_size=0.3, random_state=42)

[4] #Step 4: Initialize and Train Models
# Initialize models
dt_model = DecisionTreeClassifier(random_state=42)
nb_model = GaussianNB()
svm_model = SVC(random_state=42)
knn_model = KNeighborsClassifier()

[5] # Train models
dt_model.fit(X_train, y_train)
nb_model.fit(X_train, y_train)
svm_model.fit(X_train, y_train)
knn_model.fit(X_train, y_train)
```



```
[6] #Step 5: Make Predictions
dt_pred = dt_model.predict(X_test)
nb_pred = nb_model.predict(X_test)
svm_pred = svm_model.predict(X_test)
knn_pred = knn_model.predict(X_test)

[7] #Step 6: Evaluate Models
#Use classification_report to get precision, recall, F1-score:

print("Decision Tree:\n", classification_report(y_test, dt_pred))
print("Naive Bayes:\n", classification_report(y_test, nb_pred))
print("SVM:\n", classification_report(y_test, svm_pred))
```

```
print("KNN:\n", classification_report(y_test, knn_pred))
print("Decision Tree Confusion Matrix:\n", confusion_matrix(y_test, dt_pred))
print("Naive Bayes Confusion Matrix:\n", confusion_matrix(y_test, nb_pred))
print("SVM Confusion Matrix:\n", confusion_matrix(y_test, svm_pred))
print("KNN Confusion Matrix:\n", confusion_matrix(y_test, knn_pred))
```

Decision Tree:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

Naive Bayes:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	0.92	0.96	13
2	0.93	1.00	0.96	13
accuracy			0.98	45
macro avg	0.98	0.97	0.97	45
weighted avg	0.98	0.98	0.98	45

SVM:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

KNN:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

Decision Tree Confusion Matrix:

```
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```

Naive Bayes Confusion Matrix:

```
[[19  0  0]
 [ 0 12  1]
 [ 0  0 13]]
```

SVM Confusion Matrix:

```
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```

KNN Confusion Matrix:

```
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```

Practical 7: Time series Forecasting using ARIMA

Step 1: Import Required Libraries

```
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import adfuller
from sklearn.metrics import mean_squared_error
import numpy as np
```


Step 2: Load Time Series Dataset

```
df = pd.read_csv('airline_passengers.csv')
df.columns = ['Month', 'Passengers'] # Rename for consistency if needed
df['Month'] = pd.to_datetime(df['Month'])
df.set_index('Month', inplace=True)
print(df.head())
```

Passengers	
Month	
1949-01-01	112
1949-02-01	118
1949-03-01	132
1949-04-01	129
1949-05-01	121

Step 3: Visualize the Time Series

```
df.plot(figsize=(10, 6), title='Monthly Airline Passengers')
plt.xlabel("Date")
plt.ylabel("Passengers")
plt.show()
```



Step 4: Check Stationarity

```
result = adfuller(df['Passengers'])  
print(f'ADF Statistic: {result[0]}')  
print(f'p-value: {result[1]}')
```

→ADF Statistic: 0.8153688792060463
p-value: 0.991880243437641


Step 5: Differencing if not stationary

```
if result[1] > 0.05:
```

```
    df_diff = df['Passengers'].diff().dropna()  
    df_diff.plot(figsize=(10, 6), title='Differenced Series')  
    plt.show()  
    data_to_model = df_diff  
    d = 1
```


```
else:
```

```
    data_to_model = df['Passengers']  
    d = 0
```



Step 6: ACF and PACF plots to decide p and q

```
plot_acf(data_to_model, lags=20)  
plot_pacf(data_to_model, lags=20)  
plt.show()
```



```
# Step 7: Fit ARIMA Model (Example: p=2, d=d, q=2)
model = ARIMA(df['Passengers'], order=(2, d, 2))
model_fit = model.fit()
print(model_fit.summary())
```

SARIMAX Results						
=====						
Dep. Variable:	Passengers	No. Observations:	144			
Model:	ARIMA(2, 1, 2)	Log Likelihood	-671.673			
Date:	Fri, 23 May 2025	AIC	1353.347			
Time:	23:01:55	BIC	1368.161			
Sample:	01-01-1949	HQIC	1359.366			
	- 12-01-1960					
Covariance Type:	opg					
=====						
	coef	std err	z	P> z	[0.025	0.975]

ar.L1	1.6850	0.020	83.061	0.000	1.645	1.725
ar.L2	-0.9549	0.017	-55.420	0.000	-0.989	-0.921
ma.L1	-1.8432	0.124	-14.852	0.000	-2.086	-1.600
ma.L2	0.9953	0.134	7.401	0.000	0.732	1.259
sigma2	665.9657	113.854	5.849	0.000	442.816	889.115
=====						
Ljung-Box (L1) (Q):	0.30	Jarque-Bera (JB):	1.84			
Prob(Q):	0.59	Prob(JB):	0.40			
Heteroskedasticity (H):	7.38	Skew:	0.27			
Prob(H) (two-sided):	0.00	Kurtosis:	3.14			
=====						

```
# Step 8: Forecast Future Values
forecast = model_fit.forecast(steps=12)
```

```
print("Forecast:\n", forecast)
```

```
# Plot original and forecast
```

```
plt.figure(figsize=(10, 6))
```

```
plt.plot(df.index, df['Passengers'], label='Original')
```

```
plt.plot(pd.date_range(df.index[-1], periods=13, freq='MS')[1:], forecast, label='Forecast', color='red')
```

```
plt.title('ARIMA Forecast')
```

```
plt.xlabel('Date')
```

```
plt.ylabel('Passengers')
```

```
plt.legend()
```

```
plt.show()
```

Forecast:

1961-01-01 439.854649

1961-02-01 465.296295

1961-03-01 500.666061

1961-04-01 535.971971

1961-05-01 561.690397

1961-06-01 571.314654

1961-07-01 562.974471

1961-08-01 539.731319

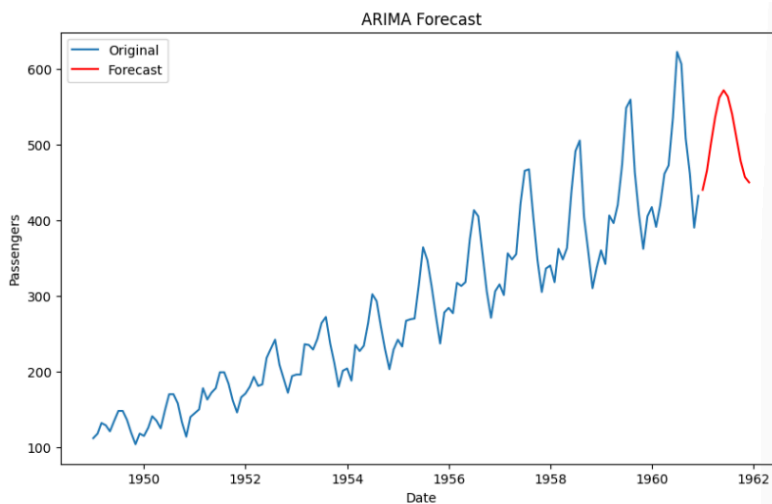
1961-09-01 508.529683

1961-10-01 478.147936

1961-11-01 456.746897

1961-12-01 449.695697

Freq: MS, Name: predicted_mean, dtype: float64



```
# Step 9: Evaluate Model Performance (Train/Test Split)
```

```
train = df['Passengers'][:-12]
```

```
test = df['Passengers'][-12:]
```

```
model = ARIMA(train, order=(2, d, 2))
```

```
model_fit = model.fit()
```

```
pred = model_fit.forecast(steps=12)
```

```
rmse = np.sqrt(mean_squared_error(test, pred))
```

```
print(f'RMSE: {rmse:.2f}') → RMSE: 55.22
```

Practical 8: Text analysis and Sentiment Detection

```
[1]import pandas as pd
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from sklearn.feature_extraction.text import TfidfVectorizer
from textblob import TextBlob
from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
# Download NLTK data (only first time)
nltk.download('punkt')
nltk.download('stopwords')
```

[2] #Step 2: Load Sample Text Data

```
#For demonstration, you can use a sample dataset of reviews, for example:
data = {'review': [
    "I loved the movie! It was fantastic and thrilling.",
    "The product is terrible, it broke after one use.",
    "Not bad, but could be better.",
    "Absolutely amazing! Highly recommend it.",
    "Worst experience ever, will not buy again."
]}
df = pd.DataFrame(data)
```

Index	Review
0	I loved the movie! It was fantastic and thrilling.
1	The product is terrible, it broke after one use.
2	Not bad, but could be better.
3	Absolutely amazing! Highly recommend it.
4	Worst experience ever, will not buy again.

[3] #Step 3: Text Preprocessing

```
#Convert to lowercase
#Tokenize text
#Remove stopwords
#(Optional) Stemming or Lemmatization
stop_words = set(stopwords.words('english'))
def preprocess_text(text):
    text = text.lower()
```



```
words = word_tokenize(text)
words = [w for w in words if w.isalpha()] #| Remove punctuation/numbers
words = [w for w in words if w not in stop_words]
return ' '.join(words)
df['cleaned_review'] = df['review'].apply(preprocess_text)
print(df[['review', 'cleaned_review']])
```

Review

I loved the movie! It was fantastic and thrilling.

The product is terrible, it broke after one use.

Not bad, but could be better.

Absolutely amazing! Highly recommend it.

Worst experience ever, will not buy again.

Cleaned Review

loved movie fantastic thrilling

product terrible broke one use

bad could better

absolutely amazing highly recommend

worst experience ever buy

[4] #Step 4: Compute TF-IDF Features

```
tfidf = TfidfVectorizer()
tfidf_matrix = tfidf.fit_transform(df['cleaned_review'])
print("TF-IDF feature names:\n", tfidf.get_feature_names_out())
print("TF-IDF matrix shape:", tfidf_matrix.shape)
```

TF-IDF feature names:

['absolutely' 'amazing' 'bad' 'better' 'broke' 'buy' 'could' 'ever'
'experience' 'fantastic' 'highly' 'loved' 'movie' 'one' 'product'
'recommend' 'terrible' 'thrilling' 'use' 'worst']

TF-IDF matrix shape: (5, 20)

[5] # Step 5: Perform Sentiment Analysis Using TextBlob:

```
from textblob import TextBlob
from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
```

```
def get_textblob_sentiment(text):
    analysis = TextBlob(text)
    # Polarity ranges from -1 (negative) to 1 (positive)
    if analysis.sentiment.polarity > 0:
        return 'Positive'
    elif analysis.sentiment.polarity == 0:
        return 'Neutral'
    else:
        return 'Negative'
```

```
df['sentiment_textblob'] = df['review'].apply(get_textblob_sentiment)
```

```
# Using VADER (better for social media/reviews):  
analyzer = SentimentIntensityAnalyzer()
```

```
def get_vader_sentiment(text):  
    score = analyzer.polarity_scores(text)  
    compound = score['compound']  
    if compound >= 0.05:  
        return 'Positive'  
    elif compound <= -0.05:  
        return 'Negative'  
    else:  
        return 'Neutral'
```

```
df['sentiment_vader'] = df['review'].apply(get_vader_sentiment)
```

```
[6] #Step 6: View Results  
print(df[['review', 'sentiment_textblob', 'sentiment_vader']])
```


```
review sentiment_textblob \  
0 I loved the movie! It was fantastic and thrill... Positive  
1 The product is terrible, it broke after one use. Negative  
2 Not bad, but could be better. Positive  
3 Absolutely amazing! Highly recommend it. Positive  
4 Worst experience ever, will not buy again. Negative
```

```
sentiment_vader  
0 Positive  
1 Negative  
2 Positive  
3 Positive  
4 Negative
```

```
[7] from wordcloud import WordCloud  
import matplotlib.pyplot as plt
```

```
def plot_cloud(docx):  
    myWC = WordCloud().generate(docx)  
    plt.figure(figsize=(20,10))  
    plt.imshow(myWC, interpolation='bilinear')  
    plt.axis('off')  
    plt.show()
```

```
joy_docx = ''.join(df['cleaned_review']) # Combine all cleaned reviews into one string  
plot_cloud(joy_docx)
```



Practical 9: Hadoop HDFS and MapReduce Word Count

spark

SparkSession - hive

SparkContext

[Spark UI](#)

Version

v3.3.2

Master

local[8]

AppName

Databricks Shell

```
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder.getOrCreate()
```

```
sc=spark.sparkContext
```

```
data=[1,2,3,4,5,6,7,8,9]
```

```
rdd=sc.parallelize(data,4)
```

```
rdd.collect()
```

```
Out[13]: [1, 2, 3, 4, 5, 6, 7, 8, 9]
```

```
rdd.glom().collect()
```

```
Out[14]: [[1, 2], [3, 4], [5, 6], [7, 8, 9]]
```

```
result=rdd.count()
```

```
rdd1=sc.parallelize([1,2,3,4,5,6])
```

```
result=rdd1.filter(lambda x:x%2 == 0)
```

```
print(result.collect())
```

```
[2, 4, 6]
```

```
rdd2=sc.parallelize([1,2,3])
```

```
result=rdd2.flatMap(lambda x:(x,x*2))
```

```
print(result.collect())
```

```
[1, 2, 2, 4, 3, 6]
```

```
rdd3=sc.parallelize([('a',1),('b',2),('a',3)])
```

```
result=rdd3.reduceByKey(lambda x,y:x+y)
print(result.collect())
[('a', 4), ('b', 2)]
```

```
rdd4=sc.parallelize([('a',1),('b',2),('a',3),('b',4)])
result=rdd4.groupByKey().mapValues(list)
print(result.collect())
[('a', [1, 3]), ('b', [2, 4])]
```

```
rdd5=sc.parallelize([('a',1),('b',2)])
rdd6=sc.parallelize([('a',3),('b',4)])
result= rdd5.join(rdd6)
print(result.collect())
[('b', (2, 4)), ('a', (1, 3))]
```

```
rdd7=sc.parallelize([1,2,3,2,1])
result=rdd7.distinct()
print(result.collect())
[1, 2, 3]
```

#rdd action

```
rdd01=sc.parallelize([1,2,3,4])
result=rdd01.collect()
print(result)
[1, 2, 3, 4]
```

```
rdd01=sc.parallelize([1,2,3,4])
result=rdd01.count()
print(result)
4
```

```
rdd03=sc.parallelize([1,2,3,4])
result=rdd03.reduce(lambda x,y:x+y)
print(result)
10
```

```
result=rdd01.first()
print(result)
1
```

```
result=rdd01.take(3)
print(result)
[1, 2, 3]
```

Practical 10: Building a spark Application with PySpark

```
spark
from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()
sc=spark.sparkContext
data=[
    ("Alice",29,"Engineering"),
    ("Bob",35,"Sales"),
    ("Charlie",40,"Engineering"),
    ("David",30,"HR"),
    ("Eva",25,"Sales")
]
columns=["name","age","department"]
df=spark.createDataFrame(data,columns)
df.display()
```

df: pyspark.sql.dataframe.DataFrame = [name: string, age: long ... 1 more field]

	A ^B C name	1 ² 3 age	A ^B C department
1	Alice	29	Engineering
2	Bob,	35	Sales
3	Charlie	40	Engineering
4	David	30	HR
5	Eva	25	Sales

5 rows

```
df.printSchema()
root
 |-- name: string (nullable = true)
 |-- age: long (nullable = true)
 |-- department: string (nullable = true)
df.show()
+-----+---+-----+
| name|age| department|
```

```
+-----+---+-----+
| Alice| 29|Engineering|
| Bob,| 35|    Sales|
|Charlie| 40|Engineering|
| David| 30|    HR|
| Eva| 25|    Sales|
+-----+---+-----+
```

```
df.select("name","age").show()
```

```
+-----+---+
| name|age|
+-----+---+
| Alice| 29|
| Bob,| 35|
|Charlie| 40|
| David| 30|
| Eva| 25|
+-----+---+
```

```
df.filter(df.age>30).show()
```

```
+-----+---+-----+
| name|age| department|
+-----+---+-----+
| Bob,| 35|    Sales|
|Charlie| 40|Engineering|
+-----+---+-----+
```

```
from pyspark.sql.functions import col
```

```
df.withColumn("age_plus",col("age")+5).show()
```

```
+-----+---+-----+-----+
| name|age| department|age_plus|
+-----+---+-----+-----+
| Alice| 29|Engineering|    34|
| Bob,| 35|    Sales|    40|
|Charlie| 40|Engineering|    45|
| David| 30|    HR|    35|
| Eva| 25|    Sales|    30|
+-----+---+-----+-----+
```

```
df.groupBy("department").avg("age").show()
```

```
+-----+-----+
| department|avg(age)|
+-----+-----+
```

```
|Engineering| 34.5|  
| Sales| 30.0|  
| HR| 30.0|
```

```
+-----+-----+
```

```
df.orderBy("age",ascending=False).show()
```

```
+-----+---+-----+
```

```
| name|age| department|
```

```
+-----+---+-----+
```

```
|Charlie| 40|Engineering|
```

```
| Bob,| 35| Sales|
```

```
| David| 30| HR|
```

```
| Alice| 29|Engineering|
```

```
| Eva| 25| Sales|
```

```
+-----+---+-----+
```

```
df.orderBy("age",decending=False).show()
```

```
+-----+---+-----+
```

```
| name|age| department|
```

```
+-----+---+-----+
```

```
| Eva| 25| Sales|
```

```
| Alice| 29|Engineering|
```

```
| David| 30| HR|
```

```
| Bob,| 35| Sales|
```

```
|Charlie| 40|Engineering|
```

```
+-----+---+-----+
```

```
df.withColumnRenamed("age","employee_age").show()
```

```
+-----+-----+-----+
```

```
| name|employee_age| department|
```

```
+-----+-----+-----+
```

```
| Alice| 29|Engineering|
```

```
| Bob,| 35| Sales|
```

```
|Charlie| 40|Engineering|
```

```
| David| 30| HR|
```

```
| Eva| 25| Sales|
```

```
+-----+-----+-----+
```

```
df.drop("department").show()
```

```
+-----+---+
```

```
| name|age|
```

```
+-----+---+
```

```
| Alice| 29|
```



```
| Bob,| 35|
|Charlie| 40|
| David| 30|
| Eva| 25|
+-----+---+

df.withColumn("age",col("age").cast("string")).printSchema()

root
|-- name: string (nullable = true)
|-- age: string (nullable = true)
|-- department: string (nullable = true)

df.describe()
Out[22]: DataFrame[summary: string, name: string, age: string, department: string]

from pyspark.sql.functions import when
df.withColumn("senior",when(col("age")>30,"Yes").otherwise("NO")).show()

+-----+---+-----+-----+
| name|age| department|senior|
+-----+---+-----+-----+
| Alice| 29|Engineering| NO|
| Bob,| 35| Sales| Yes|
|Charlie| 40|Engineering| Yes|
| David| 30| HR| NO|
| Eva| 25| Sales| NO|
+-----+---+-----+-----+

df.select("department").distinct().show()

+-----+
| department|
+-----+
|Engineering|
| Sales|
| HR|
+-----+
```